

**NANYANG TECHNOLOGICAL UNIVERSITY**

**SCHOOL OF COMPUTER SCIENCE AND  
ENGINEERING**



**AI6126 Advanced Computer Vision**

**Project 1:**

**CelebAMask Face Parsing**

<b>Zheng Weixiang</b>
<b>wzheng014@e.ntu.edu.sg</b>
<b>G2103278G</b>

<u>Document</u>	ACV Project 2 Report
<u>Name</u>	Zheng Weixiang
<u>Matriculation Number</u>	G2103278G

### Introduction

This is the report for AI6126 Advanced Computer Vision Project 1 CelebAMask Face Parsing. In this report, I will briefly describe the dataset, the task I am working on and the result I achieved.

### Dataset

CelebAMask-HQ dataset is a large-scale face image dataset that has 30,000 high-resolution face images selected from the CelebA dataset. The mask of CelebAMask were manually annotated with the size of 512x512 and 19 classes including all facial components and accessories such as skin, nose, eyes, eyebrows, ears, mouth, lip, hair, hat, eyeglass, earring, necklace, neck, and cloth.[1] In this project, we will only work on a subset of the dataset will consist of only 7000 images. We will further divide the dataset into 3 parts: training set which consists of 5000 images with 5000 corresponding masks, validation set which consists of 1000 images with 1000 corresponding masks, and a testing set which consists of 1000 images without masks. Our task is to correctly label every pixel in the image. The classes available are from 0 to 18, in total 19 classes. They are "background", "skin", "nose", "eye\_glasses", "left\_eye", "right\_eye", "left\_brow", "right\_brow", "left\_ear", "right\_ear", "mouth", "upper\_lip", "lower\_lip", "hair", "hat", "earring", "necklace", "neck", and "cloth".

### Specs of your training machine

I used the school provided environment. The specification is as follows:

Environment: SCSEGPU

Number of GPUs: 1

GPU: RTX 2080 Ti

CPU: Intel(R) Xeon(R) Gold 6240 CPU @ 2.60GHz

### Models

I've tried several models including:

1. DeepLabV3+ (CVPR'2018) [2]
  - R-50 20000 iterations
  - R-50 40000 iterations
  - R-101 20000 iterations
  - R-101 40000 iterations
2. PSANet (ECCV'2018) [2]
  - R-50 80000 iterations
  - R-101 80000 iterations
3. Swin Transformer (ICCV'2021) [2] [3]

- UperNet	Swin-T	ImageNet-1K	224x224	160k iterations
- UperNet	Swin-B	ImageNet-1K	224x224	160k iterations

For all of DeepLabV3+ models and PSANet models, I have mIoU around 70-73. With fine-tune of the model, I noticed the mIoU increases particularly when the learning rate is relatively low, and then the necklace IoU shows a sign of increasing therefore pumps up the overall mIoU. The best performance I can get using this method is 73.5972618437 on Coda leaderboard & 73.87 on SCSE GPU Server. Among all of them, the best empirical result is from Swin Transformer (ICCV'2021) [3], which has mIoU of 76.7753074062(17<sup>th</sup> entry) and that is the result from Swin-T model [3]. The overall best result on leaderboard is 77.5776134513 (19<sup>th</sup> entry) and that is the result from Swin-B model [3].

### The best result

Score of 78.76 on validation set and score of 77.57 on Coda testing set.

2022-03-23 01:22:31,684 - mmseg - INFO -

Class	IoU	Acc
background	92.92	96.44
skin	93.17	96.66
nose	88.54	93.45
eye glasses	84.6	94.16
left eye	82.29	89.85
right eye	82.4	90.43
left brow	76.12	86.03
right brow	74.84	84.45
left ear	77.97	86.79
right ear	77.4	87.01
mouth	85.62	92.95
upper lip	81.31	89.79
lower lip	83.3	90.19
hair	91.42	95.71
hat	76.51	84.8
earring	56.3	69.07
necklace	29.57	35.31
neck	83.47	91.18
cloth	78.76	85.57

2022-03-23 01:22:31,685 - mmseg - INFO - Summary:


2022-03-23 01:22:31,685 - mmseg - INFO -

aAcc	mIoU	mAcc
95.13	78.76	86.31

### Number of parameters

```
>>> print(sum(p.numel() for p in model.parameters() if p.requires_grad))  
121178846
```

### Screenshots of Coda submitted results

My Competitions Help  lan729 ▾

Here are your submissions to date (▼ indicates submission on leaderboard).

#	SCORE	FILENAME	SUBMISSION DATE	SIZE (BYTES)	STATUS	✓	
1	---	submit.zip	03/18/2022 07:27:14	10102	Failed		+
2	---	submit.zip	03/18/2022 07:30:42	10352	Failed		+
3	56.5902248732	test_mask.zip	03/18/2022 13:54:04	19513	Finished		+
4	14.3067968033	submit.zip	03/18/2022 13:54:50	20799	Finished		+
5	---	submit.zip	03/18/2022 13:55:03	21671	Failed		+
6	65.9368018412	test_mask.zip	03/18/2022 16:05:12	17438	Finished		+
7	69.4789339514	test_mask.zip	03/18/2022 16:52:22	17951	Finished		+
8	71.527060521	test_mask.zip	03/18/2022 17:56:39	19022	Finished		+
9	72.6181661124	test_mask.zip	03/19/2022 03:21:38	20539	Finished		+
10	72.1039804184	test_mask.zip	03/19/2022 06:49:58	22754	Finished		+
11	73.1790141386	test_mask.zip	03/19/2022 07:58:06	24312	Finished		+
12	73.4322296056	test_mask.zip	03/19/2022 12:45:39	30364	Finished		+
13	73.5972618437	test_mask.zip	03/21/2022 07:35:40	60712	Finished		+
14	73.4474708707	test_mask.zip	03/21/2022 08:18:45	60958	Finished		+
15	73.5556670606	test_mask.zip	03/21/2022 08:53:28	61077	Finished		+
16	73.5556670606	test_mask.zip	03/21/2022 10:05:55	62349	Finished		+
17	76.7753074062	test_mask.zip	03/22/2022 06:19:13	77476	Finished		+
18	76.70037038	test_mask.zip	03/22/2022 11:21:41	83202	Finished		+
19	77.5776134513	test_mask.zip	03/23/2022 06:23:42	98815	Finished	✓	+

### The detail of best performing model


It uses a pretrained checkpoint: `swin_base_patch4_window7_224.pth`. It extends the config of `upernet_swin_tiny_patch4_window7_512x512_160k_pretrain_224x224_1K.py`. In addition, it overwrites the backbone with a more capable model.

```
backbone=dict(  
    embed_dims=128, depths=[2, 2, 18, 2], num_heads=[4, 8, 16, 32]),  
    decode_head=dict(in_channels=[128, 256, 512, 1024], num_classes=19),  
    auxiliary_head=dict(in_channels=512, num_classes=19))
```

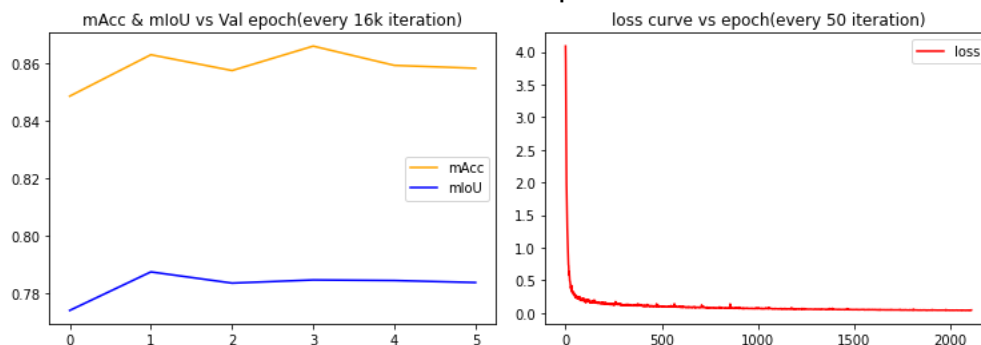
It runs 160k iterations in total, evaluate every 16k iterations, so there will be 10 evaluation epochs and 10 mIoU and mAcc in total. During the 160k iterations, log is recorded every 50 iterations. There will be 3200 entries of log recordings. Then I will show the training curves. They are plotted using `plot_curve.ipynb`, it is attached in the zip file as well.

## Training Curves

The mIoU and mAcc & Loss plot of Swin-T model



The mIoU and mAcc & Loss plot of Swin-B model



## Saved checkpoints

[https://entuedu-my.sharepoint.com/:u:/r/personal/wzheng014\\_e\\_ntu\\_edu\\_sg/Documents/ACV/iter\\_12000\\_finetune\\_7387.pth?csf=1&web=1&e=hnJMgW](https://entuedu-my.sharepoint.com/:u:/r/personal/wzheng014_e_ntu_edu_sg/Documents/ACV/iter_12000_finetune_7387.pth?csf=1&web=1&e=hnJMgW)

this is for the 15<sup>th</sup> entry 73.5556670606. It uses config file baseline.py. (attached in zip)

[https://entuedu-my.sharepoint.com/:u:/r/personal/wzheng014\\_e\\_ntu\\_edu\\_sg/Documents/ACV/iter\\_20000\\_finetune\\_7384.pth?csf=1&web=1&e=rNUINi](https://entuedu-my.sharepoint.com/:u:/r/personal/wzheng014_e_ntu_edu_sg/Documents/ACV/iter_20000_finetune_7384.pth?csf=1&web=1&e=rNUINi)

this is for the 13<sup>th</sup> entry 73.5972618437, DeepLabV3+ R-50 40000 iteration model, It uses config file baseline.py (attached in zip)

[https://entuedu-my.sharepoint.com/:u:/r/personal/wzheng014\\_e\\_ntu\\_edu\\_sg/Documents/ACV/iter\\_96000\\_swin\\_7829\\_palette.pth?csf=1&web=1&e=3rzNAu](https://entuedu-my.sharepoint.com/:u:/r/personal/wzheng014_e_ntu_edu_sg/Documents/ACV/iter_96000_swin_7829_palette.pth?csf=1&web=1&e=3rzNAu)

this is for the 17<sup>th</sup> entry, SWIN-T model, pretrained ImageNet-1K, 160k iterations. It uses config file(attached in zip)

upernet swin tiny patch4 window7 512x512 160k pretrain 224x224 1K.py

[https://entuedu-my.sharepoint.com/:u:/r/personal/wzheng014\\_e\\_ntu\\_edu\\_sg/Documents/ACV/iter\\_32000\\_swinb\\_7876.pth?csf=1&web=1&e=s6xghm](https://entuedu-my.sharepoint.com/:u:/r/personal/wzheng014_e_ntu_edu_sg/Documents/ACV/iter_32000_swinb_7876.pth?csf=1&web=1&e=s6xghm)

This is for the 19<sup>th</sup> entry, SWIN-B model, pretrained ImageNet-1K, 160k iterations. It uses config file `upernet_swin_base_512x512_160k.py`. (attached in zip)

### The ranking on Coda

The best rank I got was on Mar 23<sup>rd</sup> 2022, which is top 10.

The screenshot is as follows:

My Competitions Help lan729

Results				
#	User	Entries	Date of Last Entry	mIOU ▲
1	AccSrd	50	03/20/22	79.34809 (1)
2	zhangyu1999	39	03/23/22	79.20100 (2)
3	guozujin	16	03/23/22	78.74835 (3)
4	brucefx	43	03/22/22	78.54797 (4)
5	Nukaliad	18	03/22/22	78.53278 (5)
6	Katrina	19	03/22/22	78.15176 (6)
7	wilberpeng	9	03/20/22	77.97310 (7)
8	ignore_yuzu	6	03/22/22	77.72429 (8)
9	Annie	35	03/22/22	77.64446 (9)
10	lan729	19	03/23/22	77.57761 (10)
11	mxuai	11	03/18/22	77.29567 (11)
12	YgritteX	12	03/22/22	77.06663 (12)
13	calvenlim	20	03/22/22	77.02204 (13)
14	HeWhoMustNotBeNamed	19	03/22/22	76.93263 (14)
15	hieupham	18	03/23/22	76.47085 (15)
16	DelsinPTY	8	03/22/22	76.47044 (16)
17	gewenyue	5	03/18/22	76.42232 (17)
18	WangZiming	6	03/20/22	76.05832 (18)

### Reference

- [1] Switchablenorms. (n.d.). *Switchablenorms/celebamask-HQ: A large-scale face dataset for face parsing, recognition, generation and editing*. GitHub. Retrieved March 23, 2022, from <https://github.com/switchablenorms/CelebAMask-HQ>
- [2] Open-Mmlab. (n.d.). *Open-mmlab/mmdetection: Openmmlab semantic segmentation toolbox and benchmark*. GitHub. Retrieved March 23, 2022, from <https://github.com/open-mmlab/mmdetection>
- [3] Open-Mmlab. (n.d.). *Mmdetection/configs/swin at master · open-mmlab/mmdetection*. GitHub. Retrieved March 23, 2022, from <https://github.com/open-mmlab/mmdetection/tree/master/configs/swin>